

In [1]:

```
import torch
# Device configuration
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
device
```

Out[1]:

```
device(type='cpu')
```

In [2]:

```
from torchvision import datasets, transforms
from torchvision.transforms import ToTensor
from torch import nn
import torch.nn.functional as F
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
```

In [21]:

```
learning_rate = 0.001
batch_size    = 100
N_epochs      = 20
N_iter_train  = 250
N_iter_test   = 40
N_iter_check  = 5
input_size    = 10

hidden_size   = 10
```

In [30]:

```
def sequence_generator(L, batch_size, K = 1): #K=1, the model has to remember value at 2n
d position

    random_seq = np.random.randint(0, 9, (batch_size, L)) #generated random number sequence

    x = np.zeros((batch_size, L, 10)) #second dimension is 10 as we're looking at one-hot vectors
    y = np.zeros((batch_size, 10)) #output

    for i in range(batch_size):
        x[i, np.arange(L), random_seq[i]] = 1
        y[i, random_seq[i, K]] = 1

    #converting to torch
    random_seq = torch.tensor(random_seq, dtype=torch.float)
    x = torch.tensor(x, dtype=torch.int)
    y = torch.tensor(y, dtype=torch.float)

    return random_seq, x.float(), y #as input x is reqd to be float
```

In [23]:

```
class vanilla_RNN(nn.Module):
    def __init__(self, hidden_size):

        super().__init__()

        self.rnn = nn.RNN(input_size = 10,
                           hidden_size = hidden_size,
                           num_layers = 1,
                           bidirectional = False,
                           batch_first = True)
        self.fc = nn.Linear(hidden_size, 10)

    def forward(self, x):

        output, hidden = self.rnn(x)
        out = self.fc(output[:,-1,:])

        return out
```

In [24]:

```
from torch import optim
loss_func = nn.CrossEntropyLoss()
model = vanilla_RNN(hidden_size).to(device)
optimizer = optim.Adam(model.parameters(), lr = 0.001)
```

In [25]:

```

def Train_sequence(model,optimizer,loss_func,N_iter_train = 250):

    model.train() #setting the model in training mode
    #initializing the total training loss and total correct training predictions to 0
    train_loss      = 0
    train_correct   = 0 #correct predictions made

    train_length = batch_size*N_iter_train

    for i in range(N_iter_train):

        L = np.random.randint(3,10) #randomizing L

        random_seq,x,y = sequence_generator(L)
        x = x.to(device)
        y = y.to(device)

        pred = model(x) #prediction using the input data

        loss = loss_func(pred,y)

        optimizer.zero_grad() #zeroing out the gradients before backprop
        loss.backward()        #backprop from the loss
        optimizer.step()       #updating the weights

        pred = pred.cpu()
        loss = loss.cpu()

        #Adding this loss to training loss and computing correct predictions
        train_loss      += loss
        train_correct += (np.asarray(pred.argmax(axis = 1)-y.cpu().argmax(axis = 1))==0
        ).sum() #as subtraction will result in 0 for correct pred

        #Computing training accuracy
        train_loss = train_loss/(i+1)
        train_correct /= train_length #training accuracy

    return train_loss.detach().cpu().numpy(), train_correct #returning loss and accurac
y

```

In [26]:

```

def Test_sequence(model,loss_func,N_iter_test = N_iter_test):

    model.eval() #setting the model in eval/test mode

    #initializing the total test loss and total correct test predictions to 0
    test_loss    = 0
    test_correct = 0 #correct predictions made
    test_length  = batch_size*N_iter_test

    #switching off the gradient for eval
    with torch.no_grad():

        for i in range(N_iter_test):

            L = np.random.randint(3,10) #randomizing L

            random_seq,x,y = sequence_generator(L)
            x = x.to(device)
            y = y.to(device)

            pred = model(x) #prediction using the input data

            loss = loss_func(pred,y.argmax(axis = 1))

            pred = pred.cpu()
            loss = loss.cpu()

            #Adding this loss to testing loss and computing correct predictions
            test_loss    += loss
            test_correct += (np.asarray(pred.argmax(axis = 1))-y.cpu().argmax(axis = 1))
            ==0).sum() #as subtraction will result in 0 for correct pred

            #Computing prediction accuracy
            test_loss = test_loss/(i+1)
            test_correct /= test_length #prediction accuracy

    return test_loss.detach().cpu().numpy(), test_correct #returning loss and accuracy

```

In [27]:

```
#initialising the lists
train_losses = []
test_losses = []
train_accuracy = []
test_accuracy = []

#for epoch in range(1, N_epochs+1):
for epoch in range(1, N_epochs+1):
    print(epoch, "/", N_epochs)

    #train the model
    loss, accuracy = Train_sequence(model, optimizer, loss_func, N_iter_train = N_iter_train)

    train_losses.append(loss)
    train_accuracy.append(accuracy)
    print('Train loss for Epoch ', epoch, ': ', loss, ' | ', 'Train accuracy for Epoch ', epoch, ': ', accuracy)

    #test the model
    loss, accuracy = Test_sequence(model, loss_func, N_iter_test = N_iter_test)
    test_losses.append(loss)
    test_accuracy.append(accuracy)
    print('Test loss for Epoch ', epoch, ': ', loss, ' | ', 'Test accuracy for Epoch ', epoch, ': ', accuracy)

#Plotting the Loss and the accuracy curves
fig, (ax1, ax2) = plt.subplots(1, 2)
ax1.plot(np.asarray(train_losses), label = 'Train Loss') #converting to float array
ax1.plot(np.asarray(test_losses), label = 'Validation Loss')
ax1.set(xlabel='Epochs', ylabel='Normalised Loss')
ax1.set_title('Training and Validation error')

ax2.plot(np.asarray(train_accuracy), label = 'Training Accuracy')
ax2.plot(np.asarray(test_accuracy), label = 'Testing Accuracy')
ax2.set(xlabel='Epochs', ylabel='Accuracy')
ax2.set_title('Training and Validation accuracy')
ax2.legend()
fig.tight_layout()
```

1 / 20

Train loss for Epoch 1 : 2.2575984 | Train accuracy for Epoch 1 : 0.12164

Test loss for Epoch 1 : 2.2061765 | Test accuracy for Epoch 1 : 0.131

2 / 20

Train loss for Epoch 2 : 2.1297002 | Train accuracy for Epoch 2 : 0.17872

Test loss for Epoch 2 : 2.0904086 | Test accuracy for Epoch 2 : 0.20575

3 / 20

Train loss for Epoch 3 : 2.0817266 | Train accuracy for Epoch 3 : 0.20388

Test loss for Epoch 3 : 2.0877497 | Test accuracy for Epoch 3 : 0.199

4 / 20

Train loss for Epoch 4 : 2.0503883 | Train accuracy for Epoch 4 : 0.2298

Test loss for Epoch 4 : 2.0093217 | Test accuracy for Epoch 4 : 0.24225

5 / 20

Train loss for Epoch 5 : 1.9422538 | Train accuracy for Epoch 5 : 0.272

Test loss for Epoch 5 : 1.847606 | Test accuracy for Epoch 5 : 0.30725

6 / 20

Train loss for Epoch 6 : 1.6689517 | Train accuracy for Epoch 6 : 0.34336

Test loss for Epoch 6 : 1.5011898 | Test accuracy for Epoch 6 : 0.38525

7 / 20

Train loss for Epoch 7 : 1.3966227 | Train accuracy for Epoch 7 : 0.424

Test loss for Epoch 7 : 1.2967812 | Test accuracy for Epoch 7 : 0.458

8 / 20

Train loss for Epoch 8 : 1.213383 | Train accuracy for Epoch 8 : 0.49148

Test loss for Epoch 8 : 1.1543994 | Test accuracy for Epoch 8 : 0.52625

9 / 20

Train loss for Epoch 9 : 1.0568869 | Train accuracy for Epoch 9 : 0.5812

Test loss for Epoch 9 : 0.99806774 | Test accuracy for Epoch 9 : 0.6425

10 / 20

Train loss for Epoch 10 : 0.9453788 | Train accuracy for Epoch 10 : 0.6408

Test loss for Epoch 10 : 0.9335408 | Test accuracy for Epoch 10 : 0.6575

11 / 20

Train loss for Epoch 11 : 0.8235563 | Train accuracy for Epoch 11 : 0.68824

Test loss for Epoch 11 : 0.82683456 | Test accuracy for Epoch 11 : 0.69225

12 / 20

Train loss for Epoch 12 : 0.7469852 | Train accuracy for Epoch 12 : 0.725

Test loss for Epoch 12 : 0.63433844 | Test accuracy for Epoch 12 : 0.76125

13 / 20

Train loss for Epoch 13 : 0.6672661 | Train accuracy for Epoch 13 :

0.7534

Test loss for Epoch 13 : 0.5742372 | Test accuracy for Epoch 13 : 0.773

14 / 20

Train loss for Epoch 14 : 0.6066345 | Train accuracy for Epoch 14 : 0.77472

Test loss for Epoch 14 : 0.5709286 | Test accuracy for Epoch 14 : 0.78625

15 / 20

Train loss for Epoch 15 : 0.56055045 | Train accuracy for Epoch 15 : 0.8018

Test loss for Epoch 15 : 0.53982866 | Test accuracy for Epoch 15 : 0.808

16 / 20

Train loss for Epoch 16 : 0.48291773 | Train accuracy for Epoch 16 : 0.8354

Test loss for Epoch 16 : 0.47338468 | Test accuracy for Epoch 16 : 0.8715

17 / 20

Train loss for Epoch 17 : 0.39103648 | Train accuracy for Epoch 17 : 0.92376

Test loss for Epoch 17 : 0.37058526 | Test accuracy for Epoch 17 : 0.94375

18 / 20

Train loss for Epoch 18 : 0.30331954 | Train accuracy for Epoch 18 : 0.95852

Test loss for Epoch 18 : 0.2590236 | Test accuracy for Epoch 18 : 0.979

19 / 20

Train loss for Epoch 19 : 0.24369499 | Train accuracy for Epoch 19 : 0.97636

Test loss for Epoch 19 : 0.24138863 | Test accuracy for Epoch 19 : 0.98075

20 / 20

Train loss for Epoch 20 : 0.21773176 | Train accuracy for Epoch 20 : 0.98348

Test loss for Epoch 20 : 0.22008765 | Test accuracy for Epoch 20 : 0.983



In [33]:

```

def Check_sequence(model,lossfn,N_iter_check):

    model.eval()

    test_accuracies = []

    for L in range(3,10): #iterating through L in the required range

        #initializing the total test loss and total correct test predictions to 0
        test_loss = 0
        test_correct = 0 #correct predictions made
        test_length = N_iter_check

        #switching off the gradient for eval
        with torch.no_grad():

            random_seq,x,y = sequence_generator(L,N_iter_check)
            x = x.to(device)
            y = y.to(device)

            pred = model(x)

            loss = loss_func(pred,y.argmax(axis = 1))

            pred = pred.cpu()
            loss = loss.cpu()

            #Adding this loss to testing loss and computing correct predictions
            test_loss += loss
            test_correct += (np.asarray(pred.argmax(axis = 1))-y.cpu().argmax(axis = 1))
            ==0).sum() #as subtraction will result in 0 for correct pred

            #Computing prediction accuracy

            test_correct /= test_length #prediction accuracy

            test_accuracies.append(test_correct)

    #plotting test accuracies vs L

    plt.bar(np.arange(3,10),test_accuracies)
    plt.xlabel('L')
    plt.ylabel('Accuracy')
    plt.ylim(0,1) #as accuracy is between 0 and 1
    plt.grid()
    plt.legend()
    plt.title('Prediction Accuracy vs Length')
    plt.show()

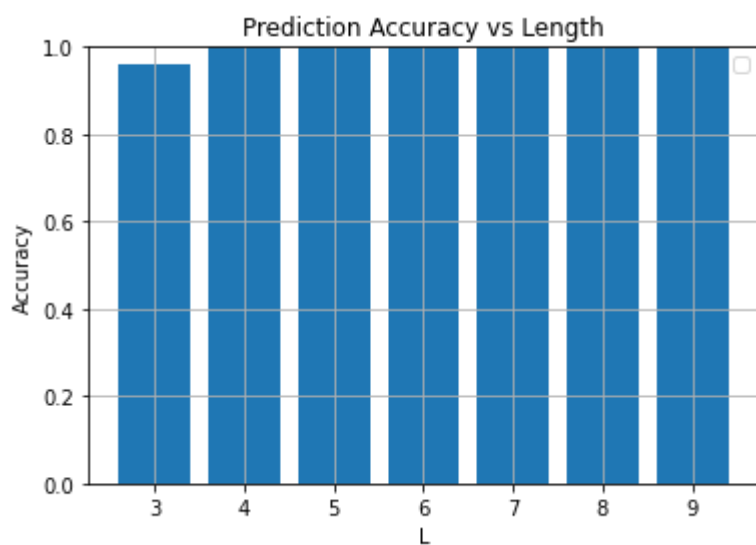
    print(f'Prediction Accuracies : {test_accuracies}')

```


In [35]:

```
Check_sequence(model, loss_func, 100)
```

WARNING:matplotlib.legend:No handles with labels found to put in legend.



Prediction Accuracies : [0.96, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0]